Shashank Anchuri

+91 8500801752 | shashankanchuri@gmail.com | LinkedIn | GitHub | shashankanchuri.space | Context Window Podcast

Education

Chaitanya Bharathi Institute of Technology (CBIT) — Hyderabad, Telangana
B.E. in Artificial Intelligence and Data Science — *Graduated May 2025*

- Coursework: Machine Learning, Deep Learning, NLP, Distributed Systems, DBMS, Advanced Algorithms.

Professional Experience

Nextenti Tech Private Limited — Hyderabad, Telangana
Software Development Engineer I — *February 2026 - Present*

- Lead a 6-engineer team (4 interns converted to full-time) building a multi-tenant healthcare ATS on Next.js, FastAPI and PostgreSQL, live with a major hospital network as the first enterprise client.
- Built AI candidate matching and ranking that scores applicants against job requirements, served over REST APIs.
- Designed tenant isolation on shared PostgreSQL with Row Level Security, enforcing access at the database layer.
- Built a platform-level AI screening chatbot as a standalone service (rules plus LLM intent routing, RAG over a healthcare corpus) serving 50+ daily sessions; tuned its MongoDB store with compound indexes.
- Added AI-driven end-to-end QA (my open-source Claude Code plugin) to releases, catching 10+ pre-release bugs.

Independent AI Consultant — Remote (Singapore-based client)
Self-Employed — *October 2025 - Present*

- Built anomaly detection over VMware error logs to flag emerging failures early, owned end to end.
- Added a RAG layer that retrieves similar past incidents and generates suggested fixes on a fully local LLM stack (Ollama, Qwen 2.5 14B), removing per-query API spend and keeping client logs on-premises.

Yellow.ai — Bangalore, Karnataka
AI Developer — *April 2025 - September 2025*

- Built LLM customer support bots and internal automation workflows orchestrating LLM calls with microservices.
- Designed reusable multi-tenant components across client deployments, cutting per-client build time.

Projects

FramesBond: AI Video Editing and Publishing Studio (in progress) | *TypeScript, Python*

- Vision-model highlight scoring that rates, tags and filters clips, with cost estimates and a spending cap per run.
- Beat, drop and energy analysis cuts highlights to music; renders long-form videos and vertical Shorts, publishes on a schedule, and runs custom detectors as sandboxed WebAssembly plugins.

Research Cartographer (GitHub) | *FastAPI, Next.js, Ollama, Sentence-Transformers*

- Local RAG pipeline: retrieves arXiv and OpenAlex papers, reranks with a cross-encoder, then a local LLM extracts paper cards and synthesises a field map (clusters, open problems), streamed to the UI over SSE.

Catch Me If You Scan (GitHub) | *Claude Code plugin, Playwright*

- Agentic QA plugin that tests web apps like a senior manual tester: derives scenarios from user stories, sweeps UI edge cases, tests multi-role flows, monitors network traffic and files evidence-backed bug reports.

Publications

A Hybrid Deep Learning Framework for Earthquake Magnitude Prediction | *ICICT 2024, Bentham Science*

- CNN-GRU model on the STEAD seismic dataset: magnitude MAE 0.30 / RMSE 0.37, beating a CNN baseline (MAE 0.47); 99.94% accuracy on earthquake vs noise detection.

Skills

Languages: Python, SQL, TypeScript, JavaScript, Java
AI/LLM: RAG (retrieval, cross-encoder reranking), Tool Calling, Agents, LangGraph, MCP, Prompt Engineering, Ollama, Sentence-Transformers, ChromaDB, HDBSCAN, PyTorch
Backend: FastAPI, Spring Boot, Next.js, REST, Celery, Redis
Data & Infra: PostgreSQL (Row Level Security), MongoDB, Docker, Kubernetes, CI/CD